

ISYE 6501 : Homework 10

3/31/2021

Table of Contents

- 14.1.1 Mean/Mode imputation
- 14.1.2 Use regression imputation
- 14.1.3 Regression with perturbation
- 15.1 Optimization

14.1.1 Mean/Mode imputation

Using data from the breast cancer data set I was able to first validate that there was missing data, determine the column containing the missing information and replace those missing values with mode and mean data.

CODE:

```
rm(list = ls())
set.seed(42)
setwd("~/Documents/ISYE6501 Intro to Analytics Modeling/FA_SP_hw10")

#Library
install.packages("tidyverse")
library(tidyverse)
install.packages("Hmisc")
library(Hmisc)

#Ingest data into a dataframe

newdata <- function() {
  datacancer <- read.table("breast-cancer-wisconsin.data.txt", header=FALSE,
                          sep="," , stringsAsFactors = FALSE)

  return(datacancer)
}

raw_cancerdf <- newdata()
summary(raw_cancerdf)

#Viewing the raw data shows V7 has the ?s
#Determine the number of rows with ?s

colSums(raw_cancerdf == '?')
```

```

#calculate the mode for all rows of column V7
Mode = function(x){
  ta = table(x)
  tam = max(ta)
  if (all(ta == tam))
    mod = NA
  else
    if(is.numeric(x))
      mod = as.numeric(names(ta)[ta == tam])
    else
      mod = names(ta)[ta == tam]
  return(mod)
}

modeinfo = Mode(raw_cancerdf$V7)
modeinfo

row_missingdata <- which(raw_cancerdf$V7 == '?', arr.ind=T)
row_missingdata

#Create a dataframe from the raw data to update ?s with the mode
mode_cancerdf <- raw_cancerdf

#Using a for statement replace all ?s with the mode of column V7

for (i in 1:nrow(mode_cancerdf)){
  if(mode_cancerdf$V7[i] == '?') {
    print('? data found')
    mode_cancerdf$V7[i] = modeinfo
  }
}

mode_cancerdf

#Use impute method to replace the ?s with the mean for column V7

mean_cancerdf <- raw_cancerdf

meanV7<- mean(as.integer(mean_cancerdf[-row_missingdata, 'V7']))
meanV7

mean_cancerdf[row_missingdata, 'V7']<- as.integer(meanV7)
mean_cancerdf[row_missingdata, 'V7']

```

```
mean_cancerdf
```

OUTPUT:

```
summary(raw_cancerdf)
```

V1	V2	V3	V4
Min. : 61634	Min. : 1.000	Min. : 1.000	Min. : 1.000
1st Qu.: 870688	1st Qu.: 2.000	1st Qu.: 1.000	1st Qu.: 1.000
Median : 1171710	Median : 4.000	Median : 1.000	Median : 1.000
Mean : 1071704	Mean : 4.418	Mean : 3.134	Mean : 3.207
3rd Qu.: 1238298	3rd Qu.: 6.000	3rd Qu.: 5.000	3rd Qu.: 5.000
Max. : 13454352	Max. : 10.000	Max. : 10.000	Max. : 10.000

V5	V6	V7	V8
Min. : 1.000	Min. : 1.000	Length:699	Min. : 1.000
1st Qu.: 1.000	1st Qu.: 2.000	Class :character	1st Qu.: 2.000
Median : 1.000	Median : 2.000	Mode :character	Median : 3.000
Mean : 2.807	Mean : 3.216		Mean : 3.438
3rd Qu.: 4.000	3rd Qu.: 4.000		3rd Qu.: 5.000
Max. : 10.000	Max. : 10.000		Max. : 10.000

V9	V10	V11
Min. : 1.000	Min. : 1.000	Min. : 2.00
1st Qu.: 1.000	1st Qu.: 1.000	1st Qu.: 2.00
Median : 1.000	Median : 1.000	Median : 2.00
Mean : 2.867	Mean : 1.589	Mean : 2.69
3rd Qu.: 4.000	3rd Qu.: 1.000	3rd Qu.: 4.00
Max. : 10.000	Max. : 10.000	Max. : 4.00

```
colSums(raw_cancerdf == '?')
```

V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11
0	0	0	0	0	0	16	0	0	0	0

```
modeinfo
```

```
[1] "1"
```

```
row_missingdata
```

```
[1] 24 41 140 146 159 165 236 250 276 293 295 298 316 322 412 618
```

After completion of the replacement of ?s with the mode the columns no longer show missing values.

```
colSums(mode_cancerdf == '?')
```

V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11
0	0	0	0	0	0	0	0	0	0	0

14.1.2 Use Regression imputation

Using the same logic, V7 is the column that is defined to have data. Linear regression is used to train the model with V7 as the response. Stepwise regression is then used to identify optimal factors for retraining the model.

CODE:

```
reg_cancerdf <- raw_cancerdf

#determine the variables that have missing data and the numbers

row_missingdata <- which(raw_cancerdf$V7 == '?', arr.ind=T)
row_missingdata

reg_cancerdf_lm <- (reg_cancerdf[-row_missingdata,2:10])
reg_cancerdf_lm$V7 <- as.integer(reg_cancerdf_lm$V7)

linear_model <- lm(V7~., data = reg_cancerdf_lm)
summary(linear_model)

#Using stepwise regression determine the optimal factors
step(linear_model)

#Using data from step process train model

linear_model2 <- lm(V7~ + V2 + V4 +V5 + V8, data = reg_cancerdf_lm)
summary(linear_model2)
```

OUTPUT:

```
row_missingdata
[1] 24 41 140 146 159 165 236 250 276 293 295 298 316 322 412 618

Call:
lm(formula = V7 ~ ., data = reg_cancerdf_lm)

Residuals:
    Min       1Q   Median       3Q      Max
-9.7316  -0.9426  -0.3002   0.6725   8.6998

Coefficients:
(Intercept)      Estimate    Std. Error  t value    Pr(>|t|)
-0.616652      0.194975    -3.163      0.00163 **
```

V2	0.230156	0.041691	5.521	4.83e-08 ***
V3	-0.067980	0.076170	-0.892	0.37246
V4	0.340442	0.073420	4.637	4.25e-06 ***
V5	0.339705	0.045919	7.398	4.13e-13 ***
V6	0.090392	0.062541	1.445	0.14883
V8	0.320577	0.059047	5.429	7.91e-08 ***
V9	0.007293	0.044486	0.164	0.86983
V10	-0.075230	0.059331	-1.268	0.20524

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.274 on 674 degrees of freedom

Multiple R-squared: 0.615, Adjusted R-squared: 0.6104

F-statistic: 134.6 on 8 and 674 DF, p-value: < 2.2e-16

step(linear_model)

Start: AIC=1131.43

V7 ~ V2 + V3 + V4 + V5 + V6 + V8 + V9 + V10

	Df	Sum of Sq	RSS	AIC
- V9	1	0.139	3486.8	1129.5
- V3	1	4.120	3490.8	1130.2
- V10	1	8.317	3495.0	1131.0
<none>			3486.6	1131.4
- V6	1	10.806	3497.5	1131.5
- V4	1	111.227	3597.9	1150.9
- V8	1	152.482	3639.1	1158.7
- V2	1	157.657	3644.3	1159.6
- V5	1	283.119	3769.8	1182.8

Step: AIC=1129.45

V7 ~ V2 + V3 + V4 + V5 + V6 + V8 + V10

	Df	Sum of Sq	RSS	AIC
- V3	1	4.028	3490.8	1128.2
- V10	1	8.179	3495.0	1129.0
<none>			3486.8	1129.5
- V6	1	11.211	3498.0	1129.7
- V4	1	114.768	3601.6	1149.6
- V2	1	158.696	3645.5	1157.8
- V8	1	160.776	3647.6	1158.2
- V5	1	285.902	3772.7	1181.3

Step: AIC=1128.24

V7 ~ V2 + V4 + V5 + V6 + V8 + V10

	Df	Sum of Sq	RSS	AIC
- V6	1	8.606	3499.4	1127.9
- V10	1	8.889	3499.7	1128.0
<none>			3490.8	1128.2
- V4	1	153.078	3643.9	1155.6
- V2	1	155.308	3646.1	1156.0
- V8	1	157.123	3647.9	1156.3
- V5	1	282.133	3772.9	1179.3

Step: AIC=1127.92

V7 ~ V2 + V4 + V5 + V8 + V10

	Df	Sum of Sq	RSS	AIC
- V10	1	5.562	3505.0	1127.0
<none>			3499.4	1127.9
- V2	1	159.594	3659.0	1156.4
- V8	1	169.954	3669.4	1158.3
- V4	1	206.785	3706.2	1165.1
- V5	1	295.807	3795.2	1181.3

Step: AIC=1127.01

V7 ~ V2 + V4 + V5 + V8

	Df	Sum of Sq	RSS	AIC
<none>			3505.0	1127.0
- V2	1	155.70	3660.7	1154.7
- V8	1	172.42	3677.4	1157.8
- V4	1	201.22	3706.2	1163.1
- V5	1	290.68	3795.7	1179.4

Call:

lm(formula = V7 ~ V2 + V4 + V5 + V8, data = reg_cancerdf_lm)

Coefficients:

(Intercept)	V2	V4	V5	V8
-0.5360	0.2262	0.3173	0.3323	0.3238

>

> #Using data from step process train model

>

> linear_model2 <- lm(V7~ + V2 + V4 +V5 + V8, data = reg_cancerdf_lm)

> summary(linear_model2)

Call:

lm(formula = V7 ~ +V2 + V4 + V5 + V8, data = reg_cancerdf_lm)

Residuals:

Min	1Q	Median	3Q	Max
-9.8115	-0.9531	-0.3111	0.6678	8.6889

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.53601	0.17514	-3.060	0.0023 **
V2	0.22617	0.04121	5.488	5.75e-08 ***
V4	0.31729	0.05086	6.239	7.76e-10 ***
V5	0.33227	0.04431	7.499	2.03e-13 ***
V8	0.32378	0.05606	5.775	1.17e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.274 on 678 degrees of freedom

Multiple R-squared: 0.6129, Adjusted R-squared: 0.6107

F-statistic: 268.4 on 4 and 678 DF, p-value: < 2.2e-16

Based on the new model the missing values are now imputed using regression and displayed.

CODE:

```
impute_model <- predict(linear_model2, reg_cancerdf[row_missingdata,])
impute_model
```

OUTPUT:

```
impute_model
  24    41   140   146   159   165   236   250
5.4585352 7.9816106 0.9872832 1.6218560 0.9807851 2.2157441 2.7152652
1.7634059
  276   293   295   298   316   322   412   618
2.0741942 6.0866099 0.9872832 2.5265324 5.2438347 1.7634059 0.9872832
0.6634986
```

14.1.3 Regression with perturbation

CODE:

```
cancerdf.perturbation <- raw_cancerdf
cancerdf.perturbation [row_missingdata, 'V7'] <-
rnorm(length(impute_model), impute_model, sd(impute_model))

cancerdf.perturbation$V7 <- as.integer(cancerdf.perturbation$V7)
```

```
cancerdf.perturbation$V7  
OUTPUT:
```

```
cancerdf.perturbation$V7  
[1] 1 10 2 4 1 10 10 1 1 1 1 1 3 3 9 1 1 1 10 1 10 7 1 8 1  
[26] 7 1 1 1 1 1 1 5 1 1 1 1 1 10 7 6 3 10 1 1 1 9 1 1 8  
[51] 3 4 5 8 8 5 6 1 10 2 3 2 8 2 1 2 1 10 9 1 1 2 1 10 4  
[76] 2 1 1 3 1 1 1 1 2 9 4 8 10 1 1 1 1 1 1 1 1 1 1 6 10  
[101] 5 5 1 3 1 3 10 10 1 9 2 9 10 8 3 5 2 10 3 2 1 2 10 10 7  
[126] 1 10 1 10 1 1 1 10 1 1 2 1 1 1 1 1 1 5 5 1 3 8 2 1 10  
[151] 1 10 5 3 1 10 1 1 1 10 10 1 1 3 1 2 10 1 1 1 1 1 1 10 10  
[176] 10 1 1 1 10 1 1 1 10 10 1 8 10 8 1 8 10 1 1 1 1 7 1 1 1  
[201] 10 10 1 1 1 10 5 1 1 1 10 8 1 10 10 5 1 1 4 1 1 10 5 8 10  
[226] 1 10 5 1 10 7 8 1 10 1 6 10 2 9 10 2 1 1 5 1 2 10 9 1 1  
[251] 1 10 10 10 8 10 1 1 1 8 10 10 10 10 3 1 10 10 4 1 10 1 10 4 1  
[276] 6 1 1 1 7 1 1 10 10 10 10 10 1 5 10 1 1 5 10 3 10 5 7 1 10  
[301] 4 1 10 1 10 10 1 1 3 5 1 1 1 1 1 2 10 8 1 5 10 1 1 10 1  
[326] 1 10 1 4 10 8 1 1 10 10 1 10 1 1 10 10 1 1 1 10 1 1 1 1 8  
[351] 1 1 3 10 1 1 3 10 4 7 10 10 3 3 1 1 10 10 1 1 1 1 1 1 1  
[376] 1 1 1 1 1 1 10 1 1 1 1 10 1 1 2 1 10 1 1 1 1 1 1 1 1  
[401] 9 1 1 4 1 1 1 1 2 1 1 0 4 1 10 3 10 1 2 1 3 10 1 1 1  
[426] 10 1 2 1 1 1 1 1 1 8 10 1 1 1 1 10 4 3 2 1 1 1 1 1 10  
[451] 1 1 1 10 1 6 10 3 1 1 1 5 1 1 1 4 10 10 1 1 1 1 1 1 1  
[476] 1 1 1 1 10 1 1 5 10 1 3 1 10 3 4 1 10 1 10 5 1 1 1 1 1  
[501] 1 1 1 1 1 1 5 4 1 1 1 1 1 1 1 10 10 1 1 1 10 1 1 5 10 1  
[526] 1 1 1 1 1 10 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 10 1 1 5  
[551] 1 1 1 5 1 1 1 1 1 1 1 1 1 1 1 1 10 1 3 10 5 10 10 1 1 2  
[576] 1 1 1 1 1 1 10 10 1 1 1 10 1 3 1 1 10 10 1 10 1 1 1 1 1  
[601] 1 1 1 1 10 8 1 1 10 1 10 2 10 1 1 1 1 2 1 1 1 2 1 1 1  
[626] 4 6 5 1 1 1 1 1 3 1 1 1 2 1 1 1 1 1 1 1 1 1 1 2 1  
[651] 4 1 1 1 1 1 1 1 10 1 1 1 1 1 1 1 1 1 1 5 8 1 1 1 1  
[676] 1 1 1 1 1 10 10 1 1 1 1 1 1 1 1 1 5 1 1 2 1 3 4 5  
>
```

Each of the methods has value in replacing the missing data. Mode/mean imputation is the easiest in terms of basic replacement but, may risk skewing values in smaller datasets.

15.1 Optimization

A good model for optimization would be determining the maximum number of classrooms needed to split students in a high school into groups of no more than 15 for standardized testing. Note that the normal class size is 28 and all classrooms are normally utilized. Variables that should be used are the number of students, the length of time allowed for testing, the number of classrooms currently available and the days that testing could be administered. The objective

function is to determine if all students can be tested on the same day or if testing needs to be split among multiple days to accommodate all students.